

Introduction to R

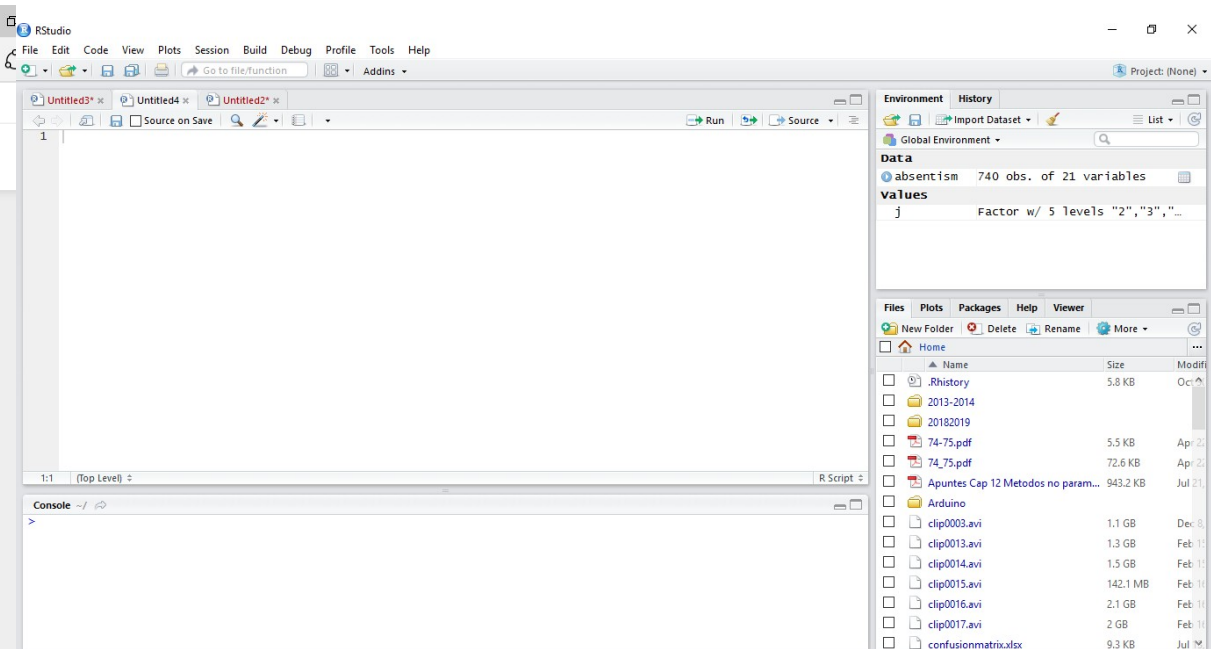
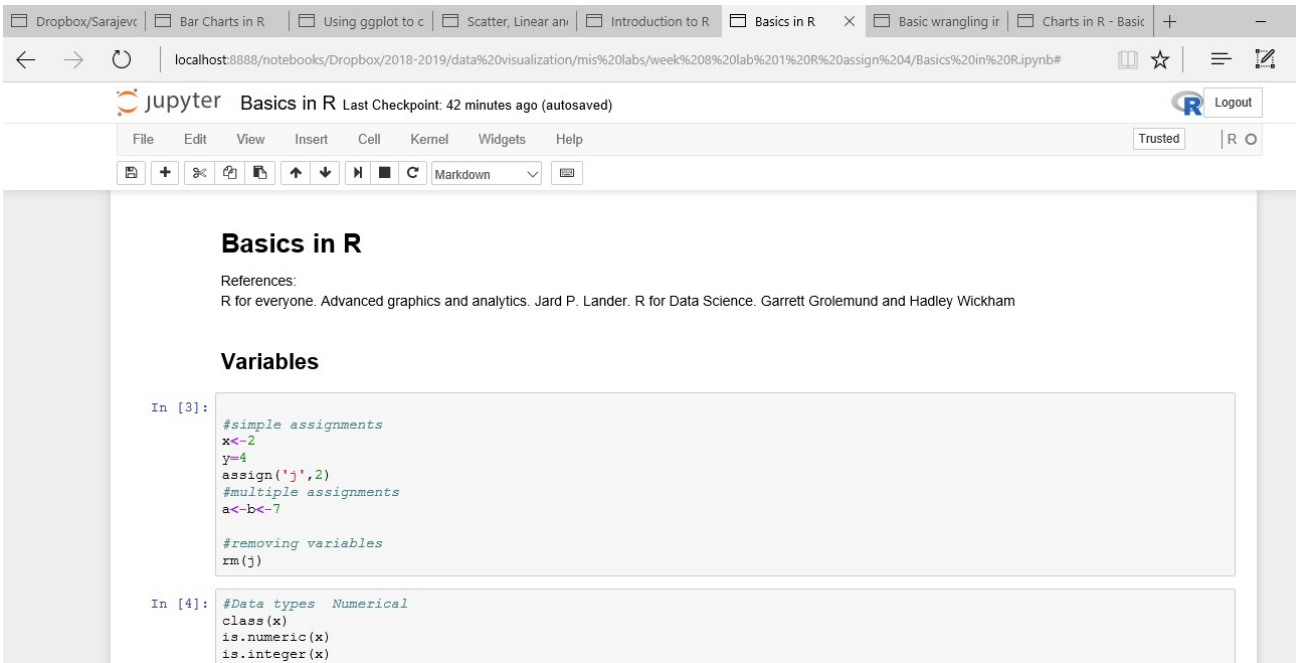
Cathy Ennis

Overview

- Introduction to R and environments
- Getting data
- Working with data
- Basic plots
- ggplot2

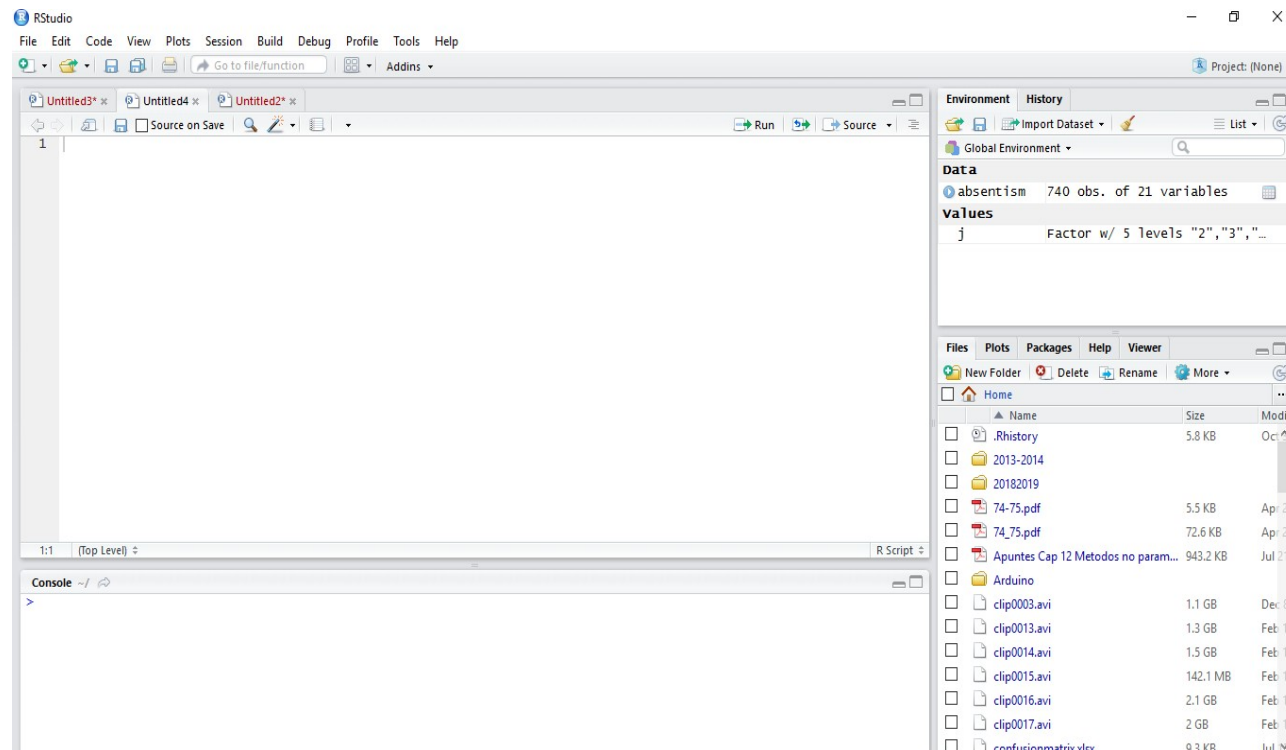
R and IDEs

- R
- Rstudio
- Jupyter Notebooks with R kernel



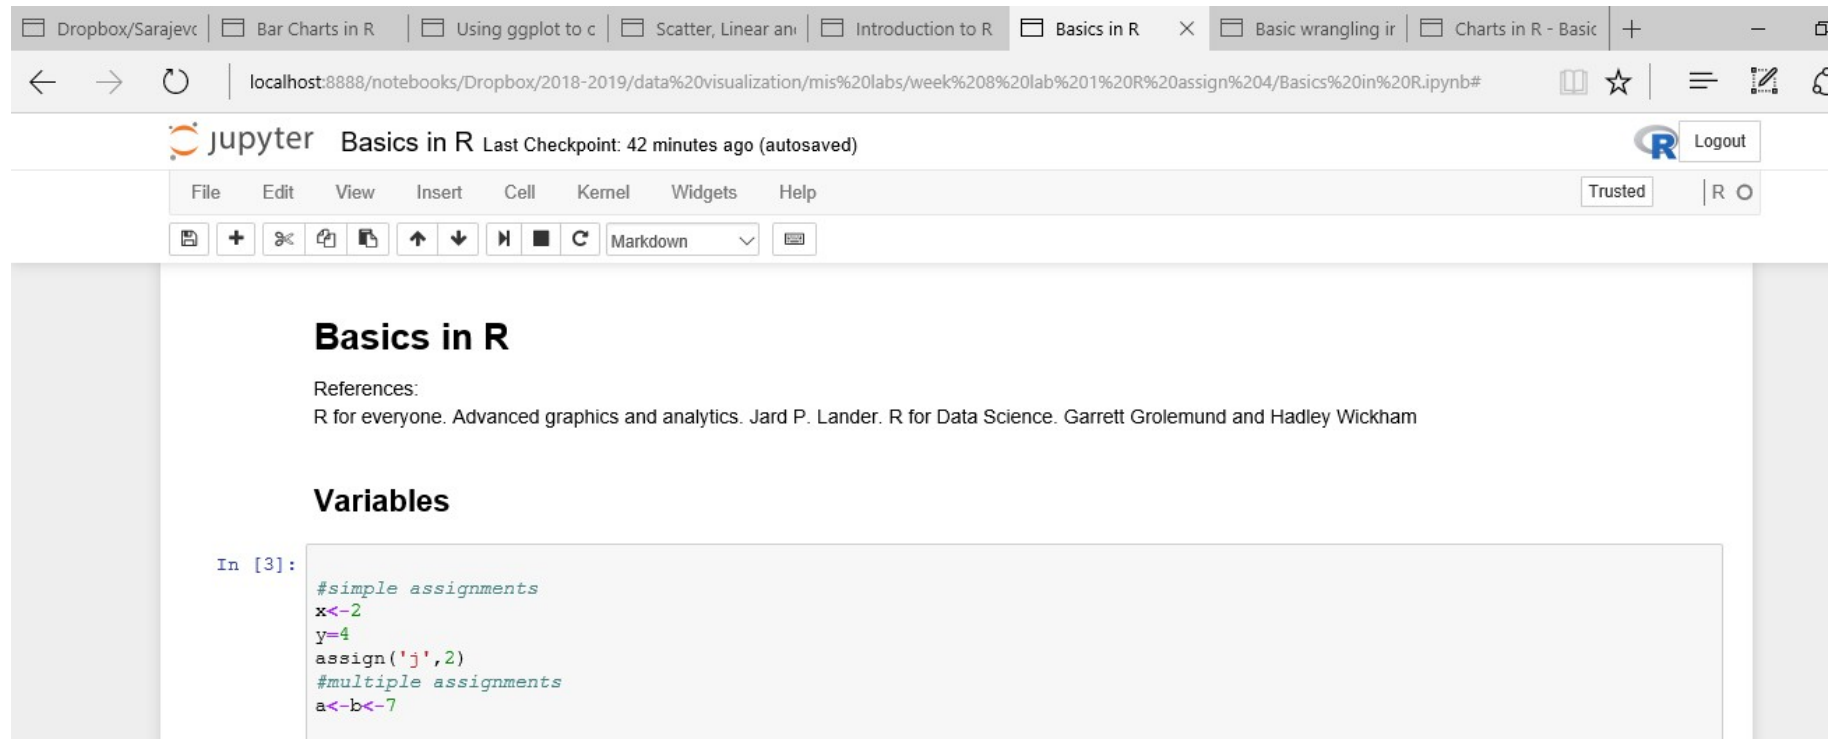
R and IDEs

- R
- Rstudio



R and IDEs

- R
- Rstudio
- Jupyter Notebooks with R kernel



The screenshot displays a web browser window with a Jupyter Notebook interface. The browser's address bar shows the URL: `localhost:8888/notebooks/Dropbox/2018-2019/data%20visualization/mis%20labs/week%208%20lab%201%20R%20assign%204/Basics%20in%20R.ipynb#`. The notebook's title bar indicates the file is named "Basics in R" and shows a "Last Checkpoint: 42 minutes ago (autosaved)". The interface includes a menu bar with options: File, Edit, View, Insert, Cell, Kernel, Widgets, and Help. Below the menu is a toolbar with icons for file operations, cell navigation, and execution. The notebook content area features a title "Basics in R" followed by a "References:" section listing "R for everyone. Advanced graphics and analytics. Jared P. Lander. R for Data Science. Garrett Golemund and Hadley Wickham". A "Variables" section is also present. The current code cell, labeled "In [3]:", contains the following R code:


```
#simple assignments
x<-2
y=4
assign('j',2)
#multiple assignments
a<-b<-7
```

Need help?

- Help for functions in R
- `?function_name` -> Description and examples
- Package documentation

Need extra help?

- Help for functions in R
 - ?function_name -> Description and examples
 - Package documentation
-
- Google it!
 - Active Community
 - Stack Overflow



Basics in R

Inspection of datasets/variables

- quickly view a dataset that has been loaded into Rstudio
- get more information about your data frame
- number of columns, number of rows, and length
- get the names of each column
- get the dimensions (number of rows and cols)

```
View(studentresult)
str(studentresult)
ncol(studentresult)
nrow(studentresult)
length(studentresult)
names(studentresult)
dim(studentresult)
```


Basics in R

The main data structures in R are:

- **list** - Contains data elements of any type
- **vector** - Sequence of data elements of the same basic type
- **matrix** - Vector with additional dimensions
- **data frame** - Used to store data tables, list of vectors of equal length

The data primitives in R are:

- **numeric** - Numeric data (approximations of the real numbers, \mathbb{R})
- **integer** - Integer data (whole numbers, \mathbb{Z})
- **factor** - Categorical data (simple classifications, like gender)
- **ordered** - Ordinal data (ordered classifications)
- **character** - Character data (strings)
- **raw** - Binary data

Basics in R

- Variable types
- Manipulating variables
- Slicing/accessing values within variables
 - `vector3<-c("a","b","ccc","dd","eeee")`
`vector3[1]`
 - `vector3[2:3]` `vector3[c(2,5)]` `vector3[-3]`
- Advanced data structures
 - Data frames
- Reading data `adult<-read.table(...`
- Vectorized calculations
 - `length(which(adult$occupation == " ?"))`

```
#simple assignments
```

```
x<-2
```

```
y=4
```

```
assign('j',2)
```

```
#multiple assignments
```

```
a<-b<-7
```

```
#removing variables
```

```
rm(j)
```

```
#Data types Numerical
```

```
class(x)
```

```
is.numeric(x)
```

```
is.character(x)
```

```
# #Data types Characters
```

```
xchar<- 'data'
```

```
nchar(xchar)
```

```
ychar<-factor("Data")
```

```
# #dates
```

```
date1<-as.Date("2018-09-09")
```

```
class(date1)
```

```
date2<-as.Date("2018-09-02")
```

```
date1>date2
```

Basics in R

- Variable types
- Manipulating variables
- Slicing/accessing values within variables
 - `vector3<-c("a","b","ccc","dd","eeee")`
`vector3[1]`
 - `vector3[2:3]` `vector3[c(2,5)]` `vector3[-3]`
- Advanced data structures
 - Data frames
- Reading data `adult<-read.table(...`
- Vectorized calculations
 - `length(which(adult$occupation == " ?"))`

```
##creating a vector using c
vector1<-c(1,2,3,4,5,5)
##creating a vector using :
(sequence of numbers)
vector2<-1:6
##length of a vector
length(vector1)
##Vector operations
vector1*3
vector1+vector2
#length of this addition should be 6
as well, test it!
#no needs for loops
#Comparing vectors to values
vector1<5
any(vector1<5)
all(vector1<5)
```

```
##using nchar with vectors
vector3<-
c("a","b","ccc","dd","eeee")
nchar(vector3)
vector3[1]
vector3[2:3]
vector3[c(2,5)]
vector3[-3]
####different to python
CAREFUL!!!!
```

Basics in R

- Variable types
- Manipulating variables
- Slicing/accessing values within variables
 - `vector3<-c("a","b","ccc","dd","eeee")`
`vector3[1]`
 - `vector3[2:3]` `vector3[c(2,5)]` `vector3[-3]`
- Advanced data structures
 - Data frames
- Reading data `adult<-read.table(...`
- Vectorized calculations
 - `length(which(adult$occupation == " ?"))`

```
##Factor vectors
```

```
g<-c("January","February","March","April","May","June","June")
```

```
g
```

```
gfactor<-as.factor(g)
```

```
levels(gfactor)
```

Basics in R

- Variable types
- Manipulating variables
- Slicing/accessing values within variables
 - `vector3<-c("a","b","ccc","dd","eeee")`
`vector3[1]`
 - `vector3[2:3]` `vector3[c(2,5)]` `vector3[-3]`
- Advanced data structures
 - Data frames
- Reading data `adult<-read.table(...`
- Vectorized calculations
 - `length(which(adult$occupation == " ?"))`

```
x<-1:10
y<-4:13

comments<-c("good", "basic", "Excellent", "good",
"basic", "good",
"basic", "Excellent", "good", "Excellent")

mydf<-data.frame(x,y,comments)

mydf

##naming our columns

mydf<-
data.frame(First=x,Second=y,Opinion=comments)

mydf

##naming our columns afterwards
colnames(mydf) <- c("First", "Second", "Third")

mydf

names(mydf)
names(mydf)[2]
head(mydf)
tail(mydf)
class(mydf)
table(mydf$Third)
```

```
## #Accessing one column
mydf$Opinion

mydf[2]

# #Accessing one row

mydf[3,]

# #accessing one element

mydf[1,3]
```

Basics in R

- Variable types
- Manipulating variables
- Slicing/accessing values within variables
 - `vector3<-c("a","b","ccc","dd","eeee")`
`vector3[1]`
 - `vector3[2:3]` `vector3[c(2,5)]` `vector3[-3]`
- Advanced data structures
 - Data frames
- Reading data `adult<-read.table(...`
- Vectorized calculations
 - `length(which(adult$occupation == " ?"))`

```
# #Reading csv files
# read.table creates a data.frame with the data. ## #From URL
mydata1<-
read.table("http://www.jaredlander.com/data/Tomato%20First.csv",header=TRUE,sep=',')
head(mydata1)

mydata2<-
read.csv2("http://www.jaredlander.com/data/Tomato%20First.csv",header=TRUE,sep=',')
head(mydata2)

adult<-read.table('https://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data',header=FALSE,sep=',')

dim(adult) colnames(adult) <-
c("age","workclass","fnlwgt","education","education_num","marital_status","occupation","relationship","race","gender","capital_gain","capital_loss","hours_per_week","native_country","income_bracket")

head(adult)
```

Basics in R

- Variable types
- Manipulating variables
- Slicing/accessing values within variables
 - `vector3<-c("a","b","ccc","dd","eeee")`
`vector3[1]`
 - `vector3[2:3]` `vector3[c(2,5)]` `vector3[-3]`
- Advanced data structures
 - Data frames
- Reading data `adult<-read.table(...`
- Vectorized calculations
 - `length(which(adult$occupation == " ?"))`

```
length(which(adult$occupation == " ?"))
any(is.na(adult$gender))
summarygender<-table(adult$gender)
summarygender[1]
summarygender[2]

aggregate(adult$age~adult$gender,adult,mean)
aggregate(adult$age~adult$gender,adult,sum)
aggregate(adult$age~adult$gender,adult,median)
?aggregate
```

Sqldf library

- Use standard SQL functions on a data frame
- sqldf allows us to run sql queries across data frames enabling us to do things like group by and sophisticated data queries.
- One of the most popular queries for cubed-type data is a group by.

To run a simple group by query in R and SQLDF, run the following command:

```
marks <- sqldf('select subject, sum(Mark_Written)  
as mark from studentresult group by subject')  
View(marks)
```

This shows the marks per student.

Without sqldf, using aggregate:

```
markswritten2<-aggregate(data=studentresult,  
Mark_Written~ Name,mean)
```


Sqldf library - Impute missing values

- Imputing missing values
- There are several ways to do this, in this example we will average scores in Oral exams over the previous two years and then use that average as the score for the missing year

The following command selects the average for Mary Healy where there are no NAs:

```
avgmark <- sqldf ( " select ROUND(AVG(Mark_Oral) )  
from studentresult where Name = 'Mary Healy' AND  
Mark_Oral is not 'NA' ")
```

Version 2 without sqldf:

```
avgmark2<- mean( studentresult$Mark_Oral [ !is.na(  
studentresult$Mark_Oral) & studentresult$Name ==  
'Mary Healy' ] )
```

Replaces those NA with the average from the query:

```
studentresult$Mark_Oral <- ifelse(  
is.na(studentresult$Mark_Oral), as.numeric(avgmark),  
studentresult$Mark_Oral)
```

Charts in R

- **ggplot**

- built on the idea that you build every graph from the same components, a data set, a set of geoms and a coordinate system.
- Layers of information can be added to customize your visualisation.

- ggplot structure

```
myplot <- ggplot(data= yourdataset, aes(x=yourx, y = youry))  
# this begins your plot by adding the data
```

Charts in R

myplot + geom_point() # this adds a geometry to your plot (scatter plot in the example)

myplot + geom_point(aes(color=dimension)) # geom layer can be customized

myplot + geom_bar()+scale_fill_brewer(palette='Reds') # customizing the colour palette

myplot + geom_bar(color=dimension) + scale_fill_manual(values=c('blue','red')) #customize colours manually

myplot+coord_map(projection="ortho", orientation= c(41,-74,0)) # map projection

myplot + theme_classic() # applies a predefined theme to the plot

myplot + labs(title="graph title", x="xaxis title", y="yaxistitle")

myplot + facet_wrap(dimension) # creates small multiples based on dimension