

Similarity Based Learning

k-Nearest Neighbour

Sarah Jane Delany

Slides adapted from ML for PDA book

Big Idea

- *Looking at what has worked well in the past and make the same (or similar) predictions*

- In 1798, Lieutenant-Colonel David Collins of HMS Calcutta was exploring in NSW when one of his sailors saw a strange animal....

| | Grrrh! |  |  | Score |
|---|--------|---|---|-------|
|  | ✓ | ✗ | ✗ | 1 |
|  | ✗ | ✓ | ✗ | 1 |
|  | ✗ | ✓ | ✓ | 2 |


Fundamentals of Similarity Learning

- **Feature space** - a D -dimensional coordinate system used to represent the descriptive features of the instances in the training data, with one axis for each feature
- **Similarity metrics** - the distance between instances in the feature space is a measure of the similarity of the instances

Feature Space

- Example: 20 college athletes and whether they were drafted by a professional team.
 - Descriptive features are Speed & Agility => 2-d feature space

| ID | Speed | Agility | Draft | ID | Speed | Agility | Draft |
|----|-------|---------|-------|----|-------|---------|-------|
| 1 | 2.50 | 6.00 | No | 11 | 2.00 | 2.00 | No |
| 2 | 3.75 | 8.00 | No | 12 | 5.00 | 2.50 | No |
| 3 | 2.25 | 5.50 | No | 13 | 8.25 | 8.50 | No |
| 4 | 3.25 | 8.25 | No | 14 | 5.75 | 8.75 | Yes |
| 5 | 2.75 | 7.50 | No | 15 | 4.75 | 6.25 | Yes |
| 6 | 4.50 | 5.00 | No | 16 | 5.50 | 6.75 | Yes |
| 7 | 3.50 | 5.25 | No | 17 | 5.25 | 9.50 | Yes |
| 8 | 3.00 | 3.25 | No | 18 | 7.00 | 4.25 | Yes |
| 9 | 4.00 | 4.00 | No | 19 | 7.50 | 8.00 | Yes |
| 10 | 4.25 | 3.75 | No | 20 | 7.25 | 5.75 | Yes |




Measuring Similarity

- Similarity between instances is measured by the distance between the instances
 - Many distance metrics - no 'best' measure => problem dependent

| ID | Speed | Agility | Draft | ID | Speed | Agility | Draft |
|----|-------|---------|-------|----|-------|---------|-------|
| 1 | 2.50 | 6.00 | No | 11 | 2.00 | 2.00 | No |
| 2 | 3.75 | 8.00 | No | 12 | 5.00 | 2.50 | No |
| 3 | 2.25 | 5.50 | No | 13 | 8.25 | 8.50 | No |
| 4 | 3.25 | 8.25 | No | 14 | 5.75 | 8.75 | Yes |
| 5 | 2.75 | 7.50 | No | 15 | 4.75 | 6.25 | Yes |
| 6 | 4.50 | 5.00 | No | 16 | 5.50 | 6.75 | Yes |
| 7 | 3.50 | 5.25 | No | 17 | 5.25 | 9.50 | Yes |
| 8 | 3.00 | 3.25 | No | 18 | 7.00 | 4.25 | Yes |
| 9 | 4.00 | 4.00 | No | 19 | 7.50 | 8.00 | Yes |
| 10 | 4.25 | 3.75 | No | 20 | 7.25 | 5.75 | Yes |

Athletes 4 and 5 are close to each other,
low distance (high similarity)

Athletes 10 and 19 are far from
each other, high distance (low similarity)



Distance Metrics

- Mathematically a metric must conform to the following four criteria:

1. **Non-negativity**: $metric(\mathbf{a}, \mathbf{b}) \geq 0$
2. **Identity**: $metric(\mathbf{a}, \mathbf{b}) = 0 \iff \mathbf{a} = \mathbf{b}$
3. **Symmetry**: $metric(\mathbf{a}, \mathbf{b}) = metric(\mathbf{b}, \mathbf{a})$
4. **Triangular Inequality**:
 $metric(\mathbf{a}, \mathbf{b}) \leq metric(\mathbf{a}, \mathbf{c}) + metric(\mathbf{b}, \mathbf{c})$

- Most common distance metric is **Euclidean distance** which computes the length of a straight line between two points, \mathbf{a} and \mathbf{b} :

$$Euclidean(\mathbf{a}, \mathbf{b}) = \sqrt{\sum_{i=1}^m (\mathbf{a}[i] - \mathbf{b}[i])^2}$$


Euclidean Distance

$$Euclidean(\mathbf{a}, \mathbf{b}) = \sqrt{\sum_{i=1}^m (\mathbf{a}[i] - \mathbf{b}[i])^2}$$

$$\begin{aligned} Euclidean(4, 5) &= \sqrt{(3.25 - 2.75)^2 + (8.25 - 7.50)^2} \\ &= \sqrt{(0.25)^2 + (0.75)^2} \\ &= \sqrt{0.625} = 0.7906 \end{aligned}$$

$$Euclidean(10, 19) = ???$$

Note: Distance and similarity have an inverse relationship




| | Speed | Agility |
|----|-------|---------|
| 4 | 3.25 | 8.25 |
| 5 | 2.75 | 7.50 |
| 10 | 4.25 | 3.75 |
| 19 | 7.50 | 8.00 |

Manhattan Distance

- Another, less well known distance measure is the **Manhattan** distance or *taxi-cab* distance

$$Manhattan(\mathbf{a}, \mathbf{b}) = \sum_{i=1}^m abs(\mathbf{a}[i] - \mathbf{b}[i])$$




Manhattan Distance

| | Speed | Agility |
|----|-------|---------|
| 4 | 3.25 | 8.25 |
| 5 | 2.75 | 7.50 |
| 10 | 4.25 | 3.75 |
| 19 | 7.50 | 8.00 |

$$\text{Manhattan}(\mathbf{a}, \mathbf{b}) = \sum_{i=1}^m \text{abs}(\mathbf{a}[i] - \mathbf{b}[i])$$

$$\begin{aligned}\text{Manhattan}(4, 5) &= \text{abs}(3.25 - 2.75) + \text{abs}(8.25 - 7.5) \\ &= 0.5 + 0.75 = 1.25\end{aligned}$$

$$\text{Manhattan}(10, 19) = ??$$



Note: Manhattan has a slight computational advantage over Euclidean

Minkowski Distance

- Euclidean and Manhattan distances are special cases of **Minkowski** distance:

$$Minkowski(\mathbf{a}, \mathbf{b}) = \left(\sum_{i=1}^m abs(\mathbf{a}[i] - \mathbf{b}[i])^p \right)^{\frac{1}{p}}$$

- Different values of parameter p result in different distance measures
 - Manhattan distance for $p = 1$
 - Euclidean distance for $p = 2$
- The larger the value of p the more emphasis is placed on features with large differences in values because these differences are raised to the power of p

What about non numeric data?

- **Binary:** Takes only two values - a boolean True/False decision
e.g. married={True,False}, test_result={Pass,Fail}
- **Categorical (Nominal):** A feature that takes values from a finite set of values, with no intrinsic ordering to the values
e.g. blood_group={A,B,AB,O}, nationality={French,Irish,Italian}
- **Ordinal:** Similar to a categorical variable, but there is a clear ordering of the variables.
e.g. grade={A,B,C,D,E,F}, dosage={Low,Medium,High}
- **Interval:** Values that allow ordering and subtraction, but do not allow other arithmetic operations
e.g date, time

Categorical Data

- Overlap Difference (feature level):**

Simplest distance measure. Returns 0 if the two values for a feature are equal and 1 otherwise

| <i>Athlete</i> | <i>Gender</i> | <i>Nationality</i> |
|----------------|---------------|--------------------|
| x1 | Female | Irish |
| x2 | Male | Irish |
| x3 | Male | Italian |

For feature
Gender

$$d_g(x1, x2) = 1$$

$$d_g(x1, x3) = 1$$

$$d_g(x2, x3) = 0$$

For feature
Nationality

$$d_n(x1, x2) = 0$$

$$d_n(x1, x3) = 1$$

$$d_n(x2, x3) = 1$$

- Hamming distance:** Distance metric for instance represented with categorical data only, = the sum of the overlap differences across all features - i.e. number of features on which two examples disagree.

$$d(x1, x2) = 1 + 0 = 1$$

$$d(x1, x3) = 1 + 1 = 2$$

$$d(x2, x3) = 0 + 1 = 1$$

Overlap distance for *Gender* +
Overlap distance for *Nationality*

Ordinal Data

- For *ordinal features*, calculate the absolute value of the difference between the two positions in the ordered list of possible values.

e.g. Ordinal Feature *Dosage*:
{Low,Medium,High} = {1, 2, 3}

$\text{diff}(\text{Low}, \text{High}) = |1-3| = 2$
 $\text{diff}(\text{Medium}, \text{Low}) = |2-1| = 1$
 $\text{diff}(\text{High}, \text{High}) = |3-3| = 0$

Heterogeneous Distance Measures

- In many datasets, the features associated with instances will have different types (e.g. continuous, categorical, ordinal etc).

- **Local distance function:** Measure the distance between two instances based on a single feature.

| <i>Athlete</i> | <i>Speed</i> | <i>Agility</i> | <i>Gender</i> | <i>Nationality</i> |
|----------------|--------------|----------------|---------------|--------------------|
| x1 | 2.50 | 6.00 | Female | Irish |
| x2 | 3.75 | 8.00 | Male | Irish |
| x3 | 2.25 | 5.50 | Male | Italian |

- e.g. distance between **x1** and **x2** in terms of *Speed*?
 - e.g. distance between **x1** and **x3** in terms of *Gender*?
 - e.g. distance between **x2** and **x3** in terms of *Nationality*?
- **Global distance function:** Measure the distance between two instances based on the combination of the local distances across all features.

Heterogeneous Distance Functions

- We can create a **global measure** from different local distance functions, using an appropriate function for each feature.

| <i>Athlete</i> | <i>Speed</i> | <i>Agility</i> | <i>Gender</i> | <i>Nationality</i> |
|----------------|--------------|----------------|---------------|--------------------|
| x1 | 2.50 | 6.00 | Female | Irish |
| x2 | 3.75 | 8.00 | Male | Irish |
| x3 | 2.25 | 5.50 | Male | Italian |

Use absolute difference for continuous features *Speed & Agility*

Use overlap for categorical features *Gender & Nationality*

$$d(x1, x2) = 1.25 + 2.0 + 1 + 0 = 4.25$$

$$d(x1, x3) = 0.25 + 0.5 + 1 + 1 = 2.75$$

$$d(x2, x3) = 1.5 + 2.5 + 0 + 1 = 5.0$$

Global distance calculated as sum over individual local distances

- Often domain expertise is required to choose an appropriate distance measure for a particular dataset.

Nearest Neighbour Algorithm (k-NN)

- Have a set of training instances and a query to be classified
- Steps:
 - Iterate across the training instances in memory and find the instance(s) that is/are the most similar (shortest distance) to the query instance in the feature space
 - Make a prediction for the query instance based on the target values of the nearest neighbour(s) of the query instance
- **1-NN** - use the most similar/closest training instance
- **k-NN** - use the k most similar/closest training instances


Example

| ID | Speed | Agility | Draft | ID | Speed | Agility | Draft |
|----|-------|---------|-------|----|-------|---------|-------|
| 1 | 2.50 | 6.00 | No | 11 | 2.00 | 2.00 | No |
| 2 | 3.75 | 8.00 | No | 12 | 5.00 | 2.50 | No |
| 3 | 2.25 | 5.50 | No | 13 | 8.25 | 8.50 | No |
| 4 | 3.25 | 8.25 | No | 14 | 5.75 | 8.75 | Yes |
| 5 | 2.75 | 7.50 | No | 15 | 4.75 | 6.25 | Yes |
| 6 | 4.50 | 5.00 | No | 16 | 5.50 | 6.75 | Yes |
| 7 | 3.50 | 5.25 | No | 17 | 5.25 | 9.50 | Yes |
| 8 | 3.00 | 3.25 | No | 18 | 7.00 | 4.25 | Yes |
| 9 | 4.00 | 4.00 | No | 19 | 7.50 | 8.00 | Yes |
| 10 | 4.25 | 3.75 | No | 20 | 7.25 | 5.75 | Yes |

- Should we draft an athlete who with the following profile?

Speed = 6.75; Agility = 3

Example




Triangle: No
Plus: Yes

Example: 1-NN

- Compute distance (using selected distance measure) of query to each training instance
 - Rank all instances based on calculated distance
 - See which is the *closest* to query
- Remember: lowest distance = highest similarity

| ID | SPEED | AGILITY | DRAFT | Dist. | ID | SPEED | AGILITY | DRAFT | Dist. |
|----|-------|---------|-------|-------|----|-------|---------|-------|-------|
| 18 | 7.00 | 4.25 | yes | 1.27 | 11 | 2.00 | 2.00 | no | 4.85 |
| 12 | 5.00 | 2.50 | no | 1.82 | 19 | 7.50 | 8.00 | yes | 5.06 |
| 10 | 4.25 | 3.75 | no | 2.61 | 3 | 2.25 | 5.50 | no | 5.15 |
| 20 | 7.25 | 5.75 | yes | 2.80 | 1 | 2.50 | 6.00 | no | 5.20 |
| 9 | 4.00 | 4.00 | no | 2.93 | 13 | 8.25 | 8.50 | no | 5.70 |
| 6 | 4.50 | 5.00 | no | 3.01 | 2 | 3.75 | 8.00 | no | 5.83 |
| 8 | 3.00 | 3.25 | no | 3.76 | 14 | 5.75 | 8.75 | yes | 5.84 |
| 15 | 4.75 | 6.25 | yes | 3.82 | 5 | 2.75 | 7.50 | no | 6.02 |
| 7 | 3.50 | 5.25 | no | 3.95 | 4 | 3.25 | 8.25 | no | 6.31 |
| 16 | 5.50 | 6.75 | yes | 3.95 | 17 | 5.25 | 9.50 | yes | 6.67 |

Example: 3-NN



Triangle: No
Plus: Yes

Example: 3-NN

- Compute distance of query to each training instance
- Rank all instances based on distance
- Choose the k nearest neighbours
- Determine predicted target by **majority vote** of target of the k nearest neighbours

| ID | SPEED | AGILITY | DRAFT | Dist. | ID | SPEED | AGILITY | DRAFT | Dist. |
|----|-------|---------|-------|-------|----|-------|---------|-------|-------|
| 18 | 7.00 | 4.25 | yes | 1.27 | 11 | 2.00 | 2.00 | no | 4.85 |
| 12 | 5.00 | 2.50 | no | 1.82 | 19 | 7.50 | 8.00 | yes | 5.06 |
| 10 | 4.25 | 3.75 | no | 2.61 | 3 | 2.25 | 5.50 | no | 5.15 |
| 20 | 7.25 | 5.75 | yes | 2.80 | 1 | 2.50 | 6.00 | no | 5.20 |
| 9 | 4.00 | 4.00 | no | 2.93 | 13 | 8.25 | 8.50 | no | 5.70 |
| 6 | 4.50 | 5.00 | no | 3.01 | 2 | 3.75 | 8.00 | no | 5.83 |
| 8 | 3.00 | 3.25 | no | 3.76 | 14 | 5.75 | 8.75 | yes | 5.84 |
| 15 | 4.75 | 6.25 | yes | 3.82 | 5 | 2.75 | 7.50 | no | 6.02 |
| 7 | 3.50 | 5.25 | no | 3.95 | 4 | 3.25 | 8.25 | no | 6.31 |
| 16 | 5.50 | 6.75 | yes | 3.95 | 17 | 5.25 | 9.50 | yes | 6.67 |

What about 4-NN?

Example: 4-NN

| ID | SPEED | AGILITY | DRAFT | Dist. | ID | SPEED | AGILITY | DRAFT | Dist. |
|----|-------|---------|-------|-------|----|-------|---------|-------|-------|
| 18 | 7.00 | 4.25 | yes | 1.27 | 11 | 2.00 | 2.00 | no | 4.85 |
| 12 | 5.00 | 2.50 | no | 1.82 | 19 | 7.50 | 8.00 | yes | 5.06 |
| 10 | 4.25 | 3.75 | no | 2.61 | 3 | 2.25 | 5.50 | no | 5.15 |
| 20 | 7.25 | 5.75 | yes | 2.80 | 1 | 2.50 | 6.00 | no | 5.20 |
| 9 | 4.00 | 4.00 | no | 2.93 | 13 | 8.25 | 8.50 | no | 5.70 |
| 6 | 4.50 | 5.00 | no | 3.01 | 2 | 3.75 | 8.00 | no | 5.83 |
| 8 | 3.00 | 3.25 | no | 3.76 | 14 | 5.75 | 8.75 | yes | 5.84 |
| 15 | 4.75 | 6.25 | yes | 3.82 | 5 | 2.75 | 7.50 | no | 6.02 |
| 7 | 3.50 | 5.25 | no | 3.95 | 4 | 3.25 | 8.25 | no | 6.31 |
| 16 | 5.50 | 6.75 | yes | 3.95 | 17 | 5.25 | 9.50 | yes | 6.67 |

- Can break ties
 - Randomly
 - Based on the sum of the nearest neighbour distances for each target class

k-NN algorithm

- *k* nearest neighbours algorithm predicts the target class with the majority vote from the set of *k* nearest neighbours to the query *q*:

$$\mathbb{M}_k(\mathbf{q}) = \operatorname{argmax}_{c \in \text{classes}(t)} \sum_{i=1}^k \delta_{t_i, c}$$

classes(*t*) is the set of target classes

t_i is the target class for instance *i*

$\delta_{i,j}$ is Kronecker's delta $\delta_{i,j} = \begin{cases} 1 & \text{when } i = j \\ 0 & \text{when } i \neq j. \end{cases}$

Weighted k -NN algorithm

- In distance weighted k Nearest Neighbour algorithm some neighbours get higher weight than others


$$\mathbb{M}_k(\mathbf{q}) = \arg \max_{c \in \text{classes}(t)} \sum_{i=1}^k w_i \times \delta_{t_i, c}$$

- Instead of a binary vote of 1 for the class of the neighbour, closest neighbours (those that are most similar to the query) get a higher weighting when deciding the prediction for the query
- Remember: similarity = inverse of distance

$$w_i = \frac{1}{\text{dist}(\mathbf{q}, \mathbf{d}_i)}$$

Tuning for k

- A simple 1-NN classifier is easy to implement. But it will be susceptible to “noise” in the data. A misclassification will occur every time a single noisy example is retrieved.
- We might decide to vary the neighbourhood size parameter k to improve the predictive performance of k -NN.
- Choosing between different settings of an algorithm is often referred to as *hyperparameter tuning* or *model selection*.
- Using a larger k (e.g. $k > 2$) can sometimes make the classifier more robust and overcome this problem.
- But when k is large ($k \rightarrow N$) and classes are *unbalanced*, we always predict the majority class.




Data Normalisation

- Consider the dataset listing salary and age information for customers and whether or not they purchased a pension plan.

| ID | Salary | Age | Purchased |
|----|--------|-----|-----------|
| 1 | 53700 | 41 | No |
| 2 | 65300 | 37 | No |
| 3 | 48900 | 45 | Yes |
| 4 | 64800 | 49 | Yes |
| 5 | 44200 | 30 | No |
| 6 | 55900 | 57 | Yes |
| 7 | 48600 | 26 | No |
| 8 | 72800 | 60 | Yes |
| 9 | 45300 | 34 | No |
| 10 | 73200 | 52 | Yes |

What should the marketing dept expect for customer aged 35, with salary of 56,000?



Data Normalisation

- Calculate distance using Euclidean distance:

| ID | Salary | Age | Purch. | Salary and Age | | Salary Only | | Age Only | |
|----|--------|-----|--------|----------------|-------|-------------|-------|----------|-------|
| | | | | Dist. | Rank. | Dist. | Rank. | Dist. | Rank. |
| 1 | 53700 | 41 | No | 2300.0078 | 2 | 2300 | 2 | 6 | 4 |
| 2 | 65300 | 37 | No | 9300.0002 | 6 | 9300 | 6 | 2 | 2 |
| 3 | 48900 | 45 | Yes | 7100.0070 | 3 | 7100 | 3 | 10 | 6 |
| 4 | 64800 | 49 | Yes | 8800.0111 | 5 | 8800 | 5 | 14 | 7 |
| 5 | 44200 | 30 | No | 11800.0011 | 8 | 11800 | 8 | 5 | 3 |
| 6 | 55900 | 57 | Yes | 102.3914 | 1 | 100 | 1 | 22 | 9 |
| 7 | 48600 | 26 | No | 7400.0055 | 4 | 7400 | 4 | 9 | 5 |
| 8 | 72800 | 60 | Yes | 16800.0186 | 9 | 16800 | 9 | 25 | 10 |
| 9 | 45300 | 34 | No | 10700.0000 | 7 | 10700 | 7 | 1 | 1 |
| 10 | 73200 | 52 | Yes | 17200.0084 | 10 | 17200 | 10 | 17 | 8 |


- Salary feature dominates the computation of distance, age feature is virtually ignored
- Due to the larger range in the Salary feature

Data Normalisation

- Use **min-max** range normalisation to rescale values to the range of [0,1]

$$a'_i = \frac{a_i - \min(a)}{\max(a) - \min(a)} \times (high - low) + low$$

$$= \frac{a_i - \min(a)}{\max(a) - \min(a)}$$



| ID | Normalized Dataset | | | Salary and Age | | Salary Only | | Age Only | |
|----|--------------------|--------|--------|----------------|------|-------------|------|----------|------|
| | Salary | Age | Purch. | Dist. | Rank | Dist. | Rank | Dist. | Rank |
| 1 | 0.3276 | 0.4412 | No | 0.1935 | 1 | 0.0793 | 2 | 0.17647 | 4 |
| 2 | 0.7276 | 0.3235 | No | 0.3260 | 2 | 0.3207 | 6 | 0.05882 | 2 |
| 3 | 0.1621 | 0.5588 | Yes | 0.3827 | 5 | 0.2448 | 3 | 0.29412 | 6 |
| 4 | 0.7103 | 0.6765 | Yes | 0.5115 | 7 | 0.3034 | 5 | 0.41176 | 7 |
| 5 | 0.0000 | 0.1176 | No | 0.4327 | 6 | 0.4069 | 8 | 0.14706 | 3 |
| 6 | 0.4034 | 0.9118 | Yes | 0.6471 | 8 | 0.0034 | 1 | 0.64706 | 9 |
| 7 | 0.1517 | 0.0000 | No | 0.3677 | 3 | 0.2552 | 4 | 0.26471 | 5 |
| 8 | 0.9862 | 1.0000 | Yes | 0.9361 | 10 | 0.5793 | 9 | 0.73529 | 10 |
| 9 | 0.0379 | 0.2353 | No | 0.3701 | 4 | 0.3690 | 7 | 0.02941 | 1 |
| 10 | 1.0000 | 0.7647 | Yes | 0.7757 | 9 | 0.5931 | 10 | 0.50000 | 8 |

Predicting Continuous Targets

- k nearest neighbours algorithm predicts the average target value of k nearest neighbours to the query q :

$$\mathbb{M}_k(\mathbf{q}) = \frac{1}{k} \sum_{i=1}^k t_i$$

t_i is the target feature value for instance i

k-NN Algorithm

- Similarity based learning attempts to mimic a very human way of reasoning - this makes them more **explainable** - easy to interpret and understand
 - Can give people more confidence in the model
- **Lazy learning technique**: Model is built at classification time rather than training time
 - Uses a set of local training instances to classify each query
 - Appropriate for heterogeneous data
 - Computationally more expensive as the number of training instances becomes larger
- Easy to add new instances to training data for re-training to handle **concept drift**
- Important to normalise data (for all prediction algorithms)